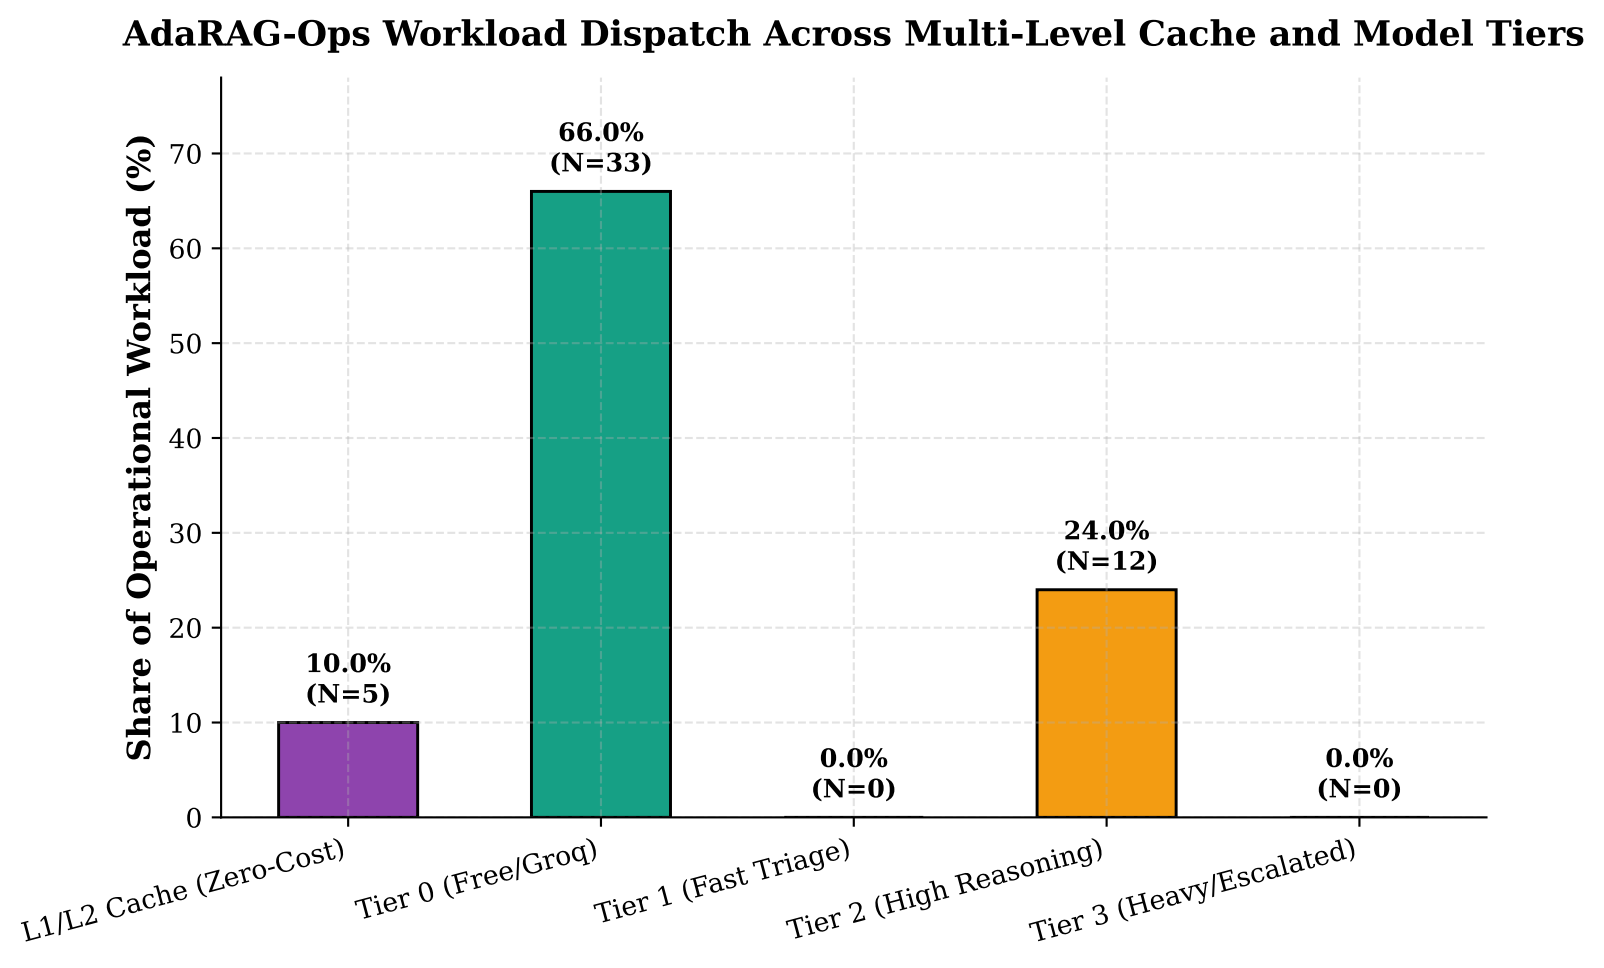
AdaRAG-Ops Workload Dispatch Across Multi-Level Cache and Model Tiers
Share of Operational Workload (%)
0
10
20
30
40
50
60
70
10.0%
(N=5)
66.0%
(N=33)
0.0%
(N=0)
24.0%
(N=12)
0.0%
(N=0)
L1/L2 Cache (Zero-Cost)
Tier 0 (Free/Groq)
Tier 1 (Fast Triage)
Tier 2 (High Reasoning)
Tier 3 (Heavy/Escalated)